

PHILLIP BRUNO

Lead Data Engineer

phillip.bruno@outlook.com | (416) 371-9919 | Toronto, ON | [LinkedIn](#) | [Portfolio](#)

PROFESSIONAL SUMMARY

Lead Data Engineer with 14+ years building and operating cloud data platforms, spanning hands-on technical leadership of AWS data infrastructure and direct people management of a six-person business intelligence team. Deep AWS specialization across ingestion, orchestration, and warehousing, with a consistent track record in data quality, observability, and pipeline reliability at 50M+ records per day. Delivers directly in Python, SQL, and Apache Spark.

TECHNICAL SKILLS

Cloud Infrastructure (AWS): Lambda, Step Functions, Glue, Redshift, S3, Athena, CloudWatch, EventBridge, EC2, IAM, CodeBuild, CodePipeline

Data Processing & Integration: Apache Spark, Airflow, PySpark, batch and streaming ingestion, CDC (Change Data Capture), API integration, event-driven architecture, ETL, ELT, data pipelines

Languages & Tools: Python (boto3, pandas, SQLAlchemy), SQL, Java, Rust, Visual Basic, Git, GitHub

Data Warehousing & BI: Redshift, Athena, BigQuery, Azure, Looker, Tableau, Power BI, QuickSight, Google Data Studio, data warehouse schema design, query optimization

Cloud Platforms: Amazon Web Services (AWS), Microsoft Azure, Google Cloud Platform (GCP)

Engineering Practices: CI/CD, Infrastructure as Code (IaC), code review, version control, data quality frameworks, data observability, data validation, anomaly detection, schema validation, data profiling

Leadership & Delivery: Technical leadership, people management, mentorship, hiring, cross-functional collaboration, roadmap planning, vendor management, stakeholder management, Agile, Scrum, sprint planning

PROFESSIONAL EXPERIENCE

Lead Data Engineer

Exostatic Inc. | June 2021 – Present

- Provide technical leadership for the data engineering function, setting code review standards, mentoring engineers, and owning the delivery roadmap for pipeline modernization.
- Architected AWS-native data ingestion pipelines using Lambda, Step Functions, and Glue, processing 50M+ records daily with 99.9% reliability and achieving 25% operational efficiency gains.
- Implemented batch and near-real-time ingestion patterns from APIs, databases, and file sources using AWS Glue jobs written in Apache Spark and PySpark alongside Lambda triggers, standardizing data onboarding across 15+ source systems.
- Developed automated data quality and observability frameworks using Lambda and Step Functions orchestration, reducing data errors by 40% and enabling real-time pipeline monitoring.
- Built CI/CD automation for data pipelines using CodeBuild and CodePipeline with GitHub as the source control backend, cutting deployment time by 60% and establishing Infrastructure as Code (IaC) practices.
- Designed Athena and Redshift data warehouse schemas optimizing query performance for analytics workloads, supporting downstream BI and reporting use cases.
- Manage vendor and partner relationships supporting the data platform, coordinating technical requirements and delivery timelines.
- Partner cross-functionally with clients in the gaming sector to launch multiple multi-million-dollar budget video games, delivering within Agile and Scrum ceremonies including sprint planning.

Sr. Manager of Business Intelligence

Omnicom Media Group | December 2019 – June 2021

- Managed a team of analysts, owning people management, performance development, and sprint planning for the business intelligence function.
- Engineered AWS, Azure, and GCP-based ETL workflows, boosting operational efficiency by 500% YoY and eliminating manual data processing.
- Built automated data validation frameworks with Python and Azure, implementing schema validation, data profiling, and anomaly detection, reducing errors by 80% for high-value clients.
- Engineered reliable analytics-ready data layers that powered executive reporting, enabling senior leadership to derive real-time strategic insights.
- Implemented CDC (Change Data Capture) techniques to monitor external teams' campaign pacing and data synchronization.
- Partnered with business owners to define data ingestion contracts and technical specifications, translating complex business rules into governed data pipelines.

Intermediate Reporting Analyst

Omnicom Media Group | December 2017 – November 2019

- Deprecated manual data processes by deploying Python-based automation scripts, achieving a 2,000-hour annual reduction in operational overhead.
- Developed reporting and data validation workflows in Python and SQL supporting client campaign analytics, establishing the ingestion and quality patterns later scaled across the BI function.

Senior Business Systems Analyst

Rogers Communications Inc. | March 2012 – November 2017

- Developed automated SQL and Java scripts to detect data anomalies and integrity issues, preventing hundreds of thousands in revenue leakage through proactive fault detection.
- Engineered a Python-based data validation framework to audit high-volume NHL log data, implementing automated quality checks that reduced error rates to under 1%.
- Developed a tool using Python and CDC techniques for validating NHL advertisements for Rogers' advertising clients, ensuring accuracy and maintaining an error rate under 0.1%.
- Partnered with business stakeholders to define data quality SLAs and technical specifications, translating complex business rules into SQL logic for automated monitoring.

EDUCATION

Bachelor of Arts (Honours): University of Toronto

Project Management Fundamentals: University of Waterloo

CERTIFICATIONS

AWS Certified Cloud Practitioner: Amazon Web Services

AWS Certified Database – Specialty: Amazon Web Services

AWARDS

Spot Award, Omnicom (2020)

REAP Award, Omnicom (2020)

Platinum Winner, Rogers (2018)

Sales Ops WOW, Rogers (2017)

Employee of the Year, Rogers (2015)